

# Pytorch 2.0

## Overall introduction

by Minh Huu Nguyen



**PYTORCH 2.X:  
FASTER,  
MORE PYTHONIC  
AND AS DYNAMIC AS EVER**

# Overall introduction

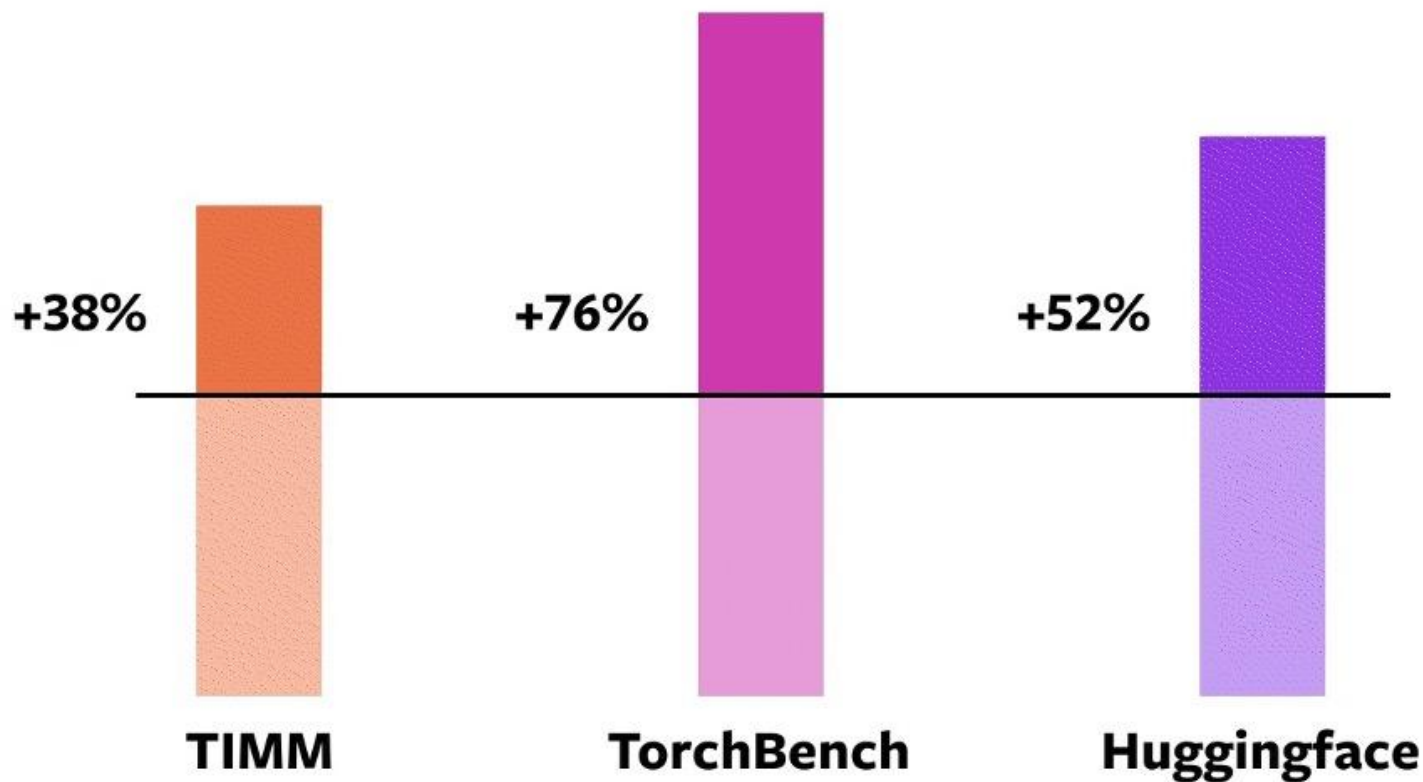
Some features of PyTorch 2.0 release are:

- OpenAI Triton deep learning compiler improve performance on Nvidia and AMD GPUs
- Accelerated Transformers with improved scaled dot product attention (SPDA)
- Metal Performance Shaders (MPS) backend accelerate PyTorch training on Mac platforms
- Amazon AWS optimizes the PyTorch CPU inference on AWS Graviton3 based C7g instances
- Some new technologies TensorParallel, DTensor, 2D parallel, TorchDynamo, AOTAutograd, PrimTorch and TorchInductor

# **torch.compile()**



# Performance



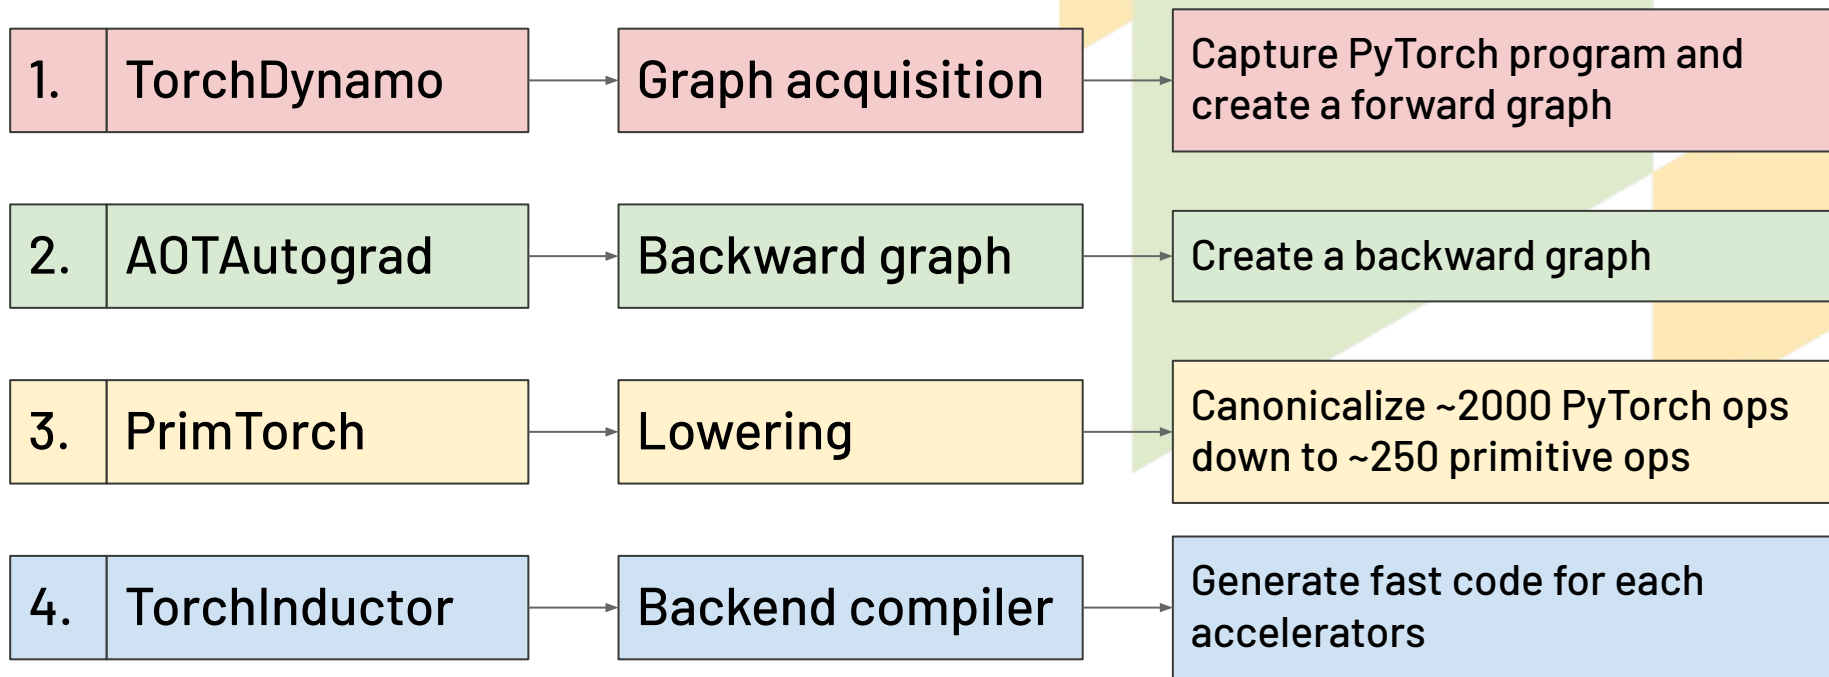
# Easy to use

```
x = torch.randn(32, 3, 64, 64)
optimizer.zero_grad()
out = model(x)
out.sum().backward()
optimizer.step()
```



```
x = torch.randn(32, 3, 64, 64)
optimizer.zero_grad()
model = torch.compile(model)
out = model(x)
out.sum().backward()
optimizer.step()
```

# Under the hood



# Easy to use

```
def torch.compile(model: Callable,
    *,
    mode: Optional[str] = "default",
    dynamic: bool = False,
    fullgraph: bool = False,
    backend: Union[str, Callable] = "inductor",
    # advanced backend options go here as kwargs
    **kwargs
) -> torch._dynamo.NNOptimizedModule
```

mode: specifies what the compiler

- default: compile efficiently
- reduce-overhead: reduce overhead by a lot more, but cost an extra memory
- max-autotune: give the fastest code, but cost a long time



